

Assignment 3: Regression Model Building Continued

Andrew G. Dunn¹

¹andrew.g.dunn@u.northwestern.edu

Andrew G. Dunn, Northwestern University Predictive Analytics Program

Prepared for PREDICT-410: Regression & Multivariate Analysis.

Formatted using markdown, pandoc, and L^AT_EX. References managed using Bibtex, and pandoc-citeproc.

Transformations – Comparisons of Y versus $\log(Y)$

Data Transformation

```
title 'Assignment 3';

libname mydata '/scs/crb519/PREDICT_410/SAS_Data/' access=readonly;

* create a temporary variable (data source is read only);
* we will also do some transformations and only keep;;
* the variables that we're interested in for this study.;
data ames;
    set mydata.ames_housing_data;
    log_SalePrice = log(SalePrice);
    log_GrLivArea = log(GrLivArea);
    keep SalePrice log_salePrice GrLivArea log_GrLivArea MasVnrArea BsmtUnfSF;

* verify that we did indeed carry over the variables of interest;
proc print data=ames (obs=5);
```

From this we see that we've retained the variables of interest for our study:

Obs	MasVnrArea	BsmtUnfSF	GrLivArea	SalePrice	log_SalePrice	log_GrLivArea
1	112	441	1656	215000	12.2784	7.41216
2	0	270	896	105000	11.5617	6.79794
3	108	406	1329	172000	12.0552	7.19218
4	0	1045	2110	244000	12.4049	7.65444
5	0	137	1629	189900	12.1543	7.39572

Table 1: Variables kept within Data Set

We'll look at the transformed variables in relation to our choice predictor variable:

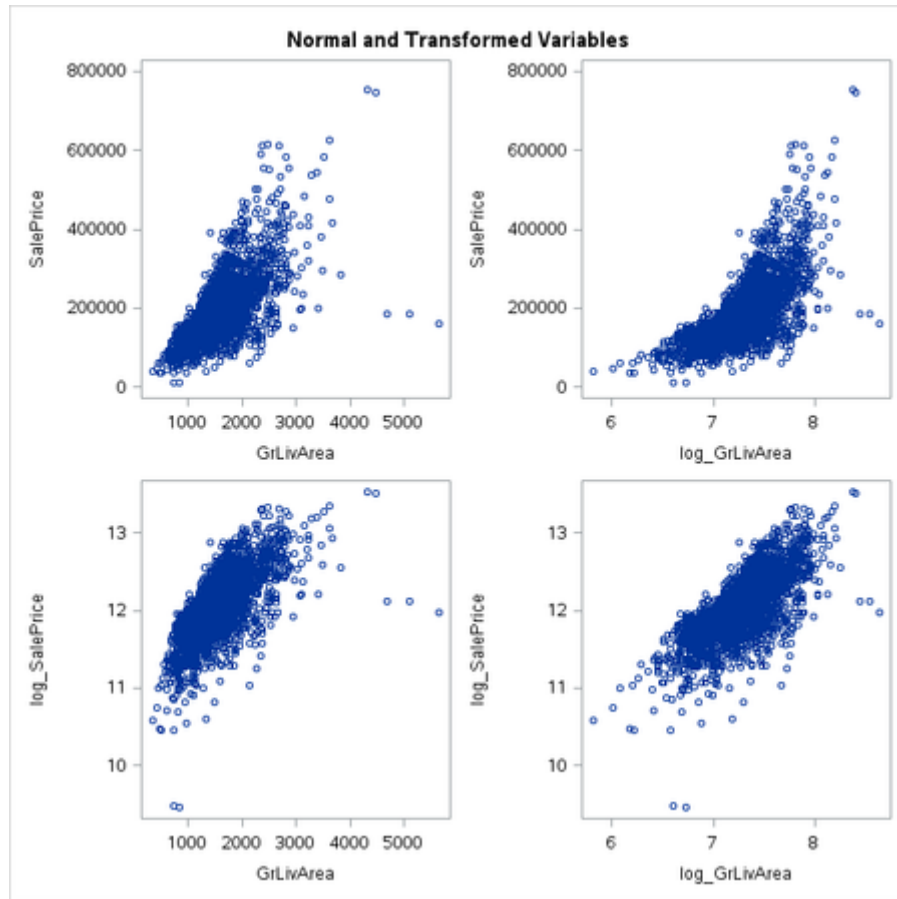


Figure 1: Normal and Transformed SalePrice, GrLivArea

It's interesting to see that of the four combinations, the combination of both scaled variables appears to be generally most linear.

Model Comparisons

We'll now compare the following models:

$$\text{SalePrice} = \beta_0 + \beta_1 \text{GrLivArea} + \epsilon$$

$$\text{SalePrice} = \beta_0 + \beta_1 \log_ \text{GrLivArea} + \epsilon$$

$$\log_ \text{SalePrice} = \beta_0 + \beta_1 \text{GrLivArea} + \epsilon$$

$$\log_ \text{SalePrice} = \beta_0 + \beta_1 \log_ \text{GrLivArea} + \epsilon$$

```
proc reg data=ames;
  model SalePrice = GrLivArea;
  model SalePrice = log_GrLivArea;
  model log_SalePrice = GrLivArea;
  model log_SalePrice = log_GrLivArea;
run;
```

Which yields the following parameters estimates:

Variable	DF	Parameter Estimate	Standard Error	t Value	Pr > t
Intercept	1	13290	3269.70277	4.06	<0.0001
GrLivArea	1	111.69400	2.06607	54.06	<0.0001

Table 2: Model: $\text{SalePrice} = 13290 + 111.694 \times \text{GrLivArea}$

As the independent variable is not transformed in this model, a one unit increase in the independent variable would result in an average change in the mean of the dependent variable by 111.694.

Variable	DF	Parameter Estimate	Standard Error	t Value	Pr > t
Intercept	1	-1060765	23758	-44.65	<0.0001
log_GrLivArea	1	171011	3269.11261	52.31	<0.0001

Table 3: Model: $\text{SalePrice} = 171011 \times \log_GrLivArea - 1060765$

As the independent variable is log-transformed in this model, a one unit increase in the independent variable would result in an average change in the mean of the dependent variable by $\frac{171011}{100}$ percent.

Variable	DF	Parameter Estimate	Standard Error	t Value	Pr > t
Intercept	1	11.17954	0.01694	660.12	<0.0001
GrLivArea	1	0.00056107	0.00001070	52.43	<0.0001

Table 4: Model: $\log_SalePrice = 11.17954 + 0.00056 \times \text{GrLivArea}$

As the dependent variable is log-transformed in this model, a one unit increase in the independent variable would result in an average change in the mean of the dependent variable by 0.00056×100 percent.

Variable	DF	Parameter Estimate	Standard Error	t Value	Pr > t
Intercept	1	5.43019	0.11644	46.63	<0.0001
log_GrLivArea	1	0.90781	0.01602	56.66	<0.0001

Table 5: Model: $\log_SalePrice = 5.43019 + 0.90781 \times \log_GrLivArea$

As both the dependent and independent variables are log-transformed in this model, a one percent increase in the independent variable would result in an average change in the mean of the dependent variable by 0.90781 percent.

We'll now aggregate the regression diagnostic information and make a simple comparison of the models.

Model	Adj R-Square	F Value
$\text{SalePrice} = 171011 \times \log_GrLivArea - 1060765$	0.4994	2922.59
$\log_SalePrice = 11.17954 + 0.00056 \times GrLivArea$	0.4829	2736.45
$\log_SalePrice = 11.17954 + 0.00056 \times GrLivArea$	0.4840	2748.89
$\log_SalePrice = 5.43019 + 0.90781 \times \log_GrLivArea$	0.5228	3209.97

Table 6: Comparison of re-expressed $\text{SalePrice} = 13290 + 111.694 \times \text{GrLivArea}$ models

If we solely look at these criteria, then we would conclude that the model that was ‘best’ based on explanation of variability in SalePrice would be the model that uses a log-transform of both SalePrice and GrLivArea. We can also look at the F Value to see that the independent and dependent variable log-transform model best first the population from which the data were sampled.

Examination of the models graphically reinforced our initial assumption that the transformation of both dependent and independent variables looked to be much more linear than alternatives. We see on the QQ-Plot of the log-transformed dependent and independent variable model that the representation is lightly tailed, but appears to be more normal than the other models.

There are still some outliers that we should be concerned about with this model fit. The model may have powerful explanatory power, however its current fit certainly incorporates some outliers from the data set.

In our reading [1] we’ve found that non-linear re-expression of the variables may be necessary when any of the following apply:

- The residuals have a skewed distribution. The purpose of a transformation is to obtain residuals that are approximately symmetrically distributed.
- The spread of the residuals changes systematically with the values of the dependent variable. The purpose of the transformation is to remove that systematic change in spread, achieving approximate homoscedasticity.
- A desire to linearize a relationship.
- When the context of the data expects, e.g. chemistry concentrations are expressed commonly as logarithms.
- A desire to simplify the model, e.g. when a log-transform can simplify the number and complexity of interaction terms.

Furthermore, a log-transform is specifically indicated instead of another transform when:

- The residuals have a “strongly” positively skewed distribution. Tukey in [2] provides quantitative ways to estimate the transformation based on rank statistics of the residuals. However, it comes down to: if a log-transform symmetrizes the residuals it was likely the right re-expression.
- When the standard deviation of the residuals is directly proportional to the fitted values.
- When the relationship is close to exponential.

Some non-reasons to use a re-expression include:

- Making outliers look less like outliers. Wuber in [1] states: an outlier is a datum that does not fit some parsimonious relatively simple description of the data. To change one’s perspective to make outliers look better is usually an incorrect reversal of priorities.

- Letting the software automatically do it (without you're desiring so).
- When making 'bad' data appear to be well behaved.
- For visualization. If you need a re-expression for visualization, ensure that you're not doing it within the model also.

Correlations to $\log(\text{SalePrice})$

```
proc corr data=ames nosimple rank;
  var log_saleprice;
  with GrLivArea MasVnrArea BsmtUnfSF BsmtFinSF1 FirstFlrSF TotalBsmtSF GarageArea;
run;
```

Variable	Pearson Correlation Coefficients	Prob > $ r $ under $H_0: \rho=0$	Number of Observations
GrLivArea	0.69586	<.0001	2930
GarageArea	0.65113	<.0001	2929
TotalBsmtSF	0.62510	<.0001	2929
FirstFlrSF	0.60263	<.0001	2930
MasVnrArea	0.44861	<.0001	2907

Table 7: Correlation of continuous variables to $\log(\text{SalePrice})$

We'll now examine how the GrLivArea variable looks with the non transformed and log-transformed SalePrice variable:

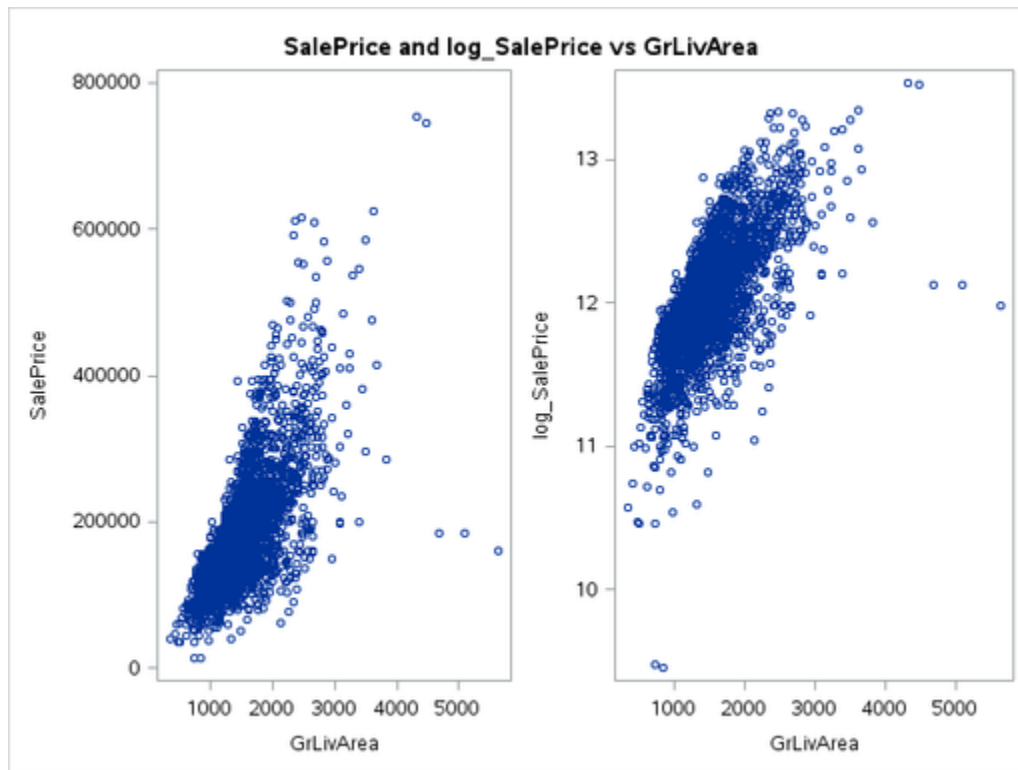


Figure 2: SalePrice and log_SalePrice vs GrLivArea

We see immediately that there is a scaling difference, which is to be expected from our transformation. If we ignored or were oblivious to the scaling, we might think that the transformed variable graph has less outliers. We refer back above where it was stated: that making bad data appear good, or doing transformations specifically for visualization purposes would not be reason enough to accept the transform.

Remarks on Transformations

We made some detailed remarks above on log-transformation in particular, transformation in general is not a strictly rule based activity, moreover it likely should be considered an art. Hubor in [3] states: One transforms the *dependent* variable to achieve approximate *symmetry* and *homoscedasticity* of the *residuals*. Where as transformations of the *independent* variable are meant to achieve *linear relationships* with the *dependent* variable.

Hubor further states that in principle there is typically nothing special about how the data are originally expressed, so one should let the data suggest re-expressions that lead to effective, accurate, useful, and theoretically justified models.

In examining the base model $\text{SalePrice} = \beta_0 + \beta_1 \text{GrLivArea} + \epsilon$ we observe that the relationship between SalePrice and GrLivArea appear to be linear. We therefor will create a transformation on the dependent variable:

$$\sqrt{\text{SalePrice}} = \beta_0 + \beta_1 \text{GrLivArea} + \epsilon$$

We will compare the models:

$$\text{SalePrice} = \beta_0 + \beta_1 \text{GrLivArea} + \epsilon$$

$$\sqrt{\text{SalePrice}} = \beta_0 + \beta_1 \text{GrLivArea} + \epsilon$$

And we'll take the best performing model from above:

$$\log_ \text{SalePrice} = 5.43019 + 0.90781 \times \log_ \text{GrLivArea}$$

Model	Adj R-Square	F Value
$\text{SalePrice} = 13290 + 111.694 \times \text{GrLivArea}$	0.4994	2922.59
$\sqrt{\text{SalePrice}} = 232.93209 + 0.12225 \times \text{GrLivArea}$	0.5073	3017.28
$\log_ \text{SalePrice} = 5.43019 + 0.90781 \times \log_ \text{GrLivArea}$	0.5228	3209.97

Table 8: Comparison of re-expressed $\text{SalePrice} = 13290 + 111.694 \times \text{GrLivArea}$ Models

We feel that our chosen transform of the dependent variable to chase approximate *symmetry* and *homoscedasticity* of the *residuals* was a good choice. It performs better than the non re-expressed model, however performs slightly worse than the log transformed model from above.

Outliers

Identify Outliers in SalePrice, Prune based on a Removal Strategy

Within SAS there is the univariate procedure, which can be used to identify outliers and extreme observations. We'll use this to examine SalePrice:

```
proc univariate normal plot data=ames;  
  var SalePrice;  
  histogram SalePrice / normal (color=red w=5);
```

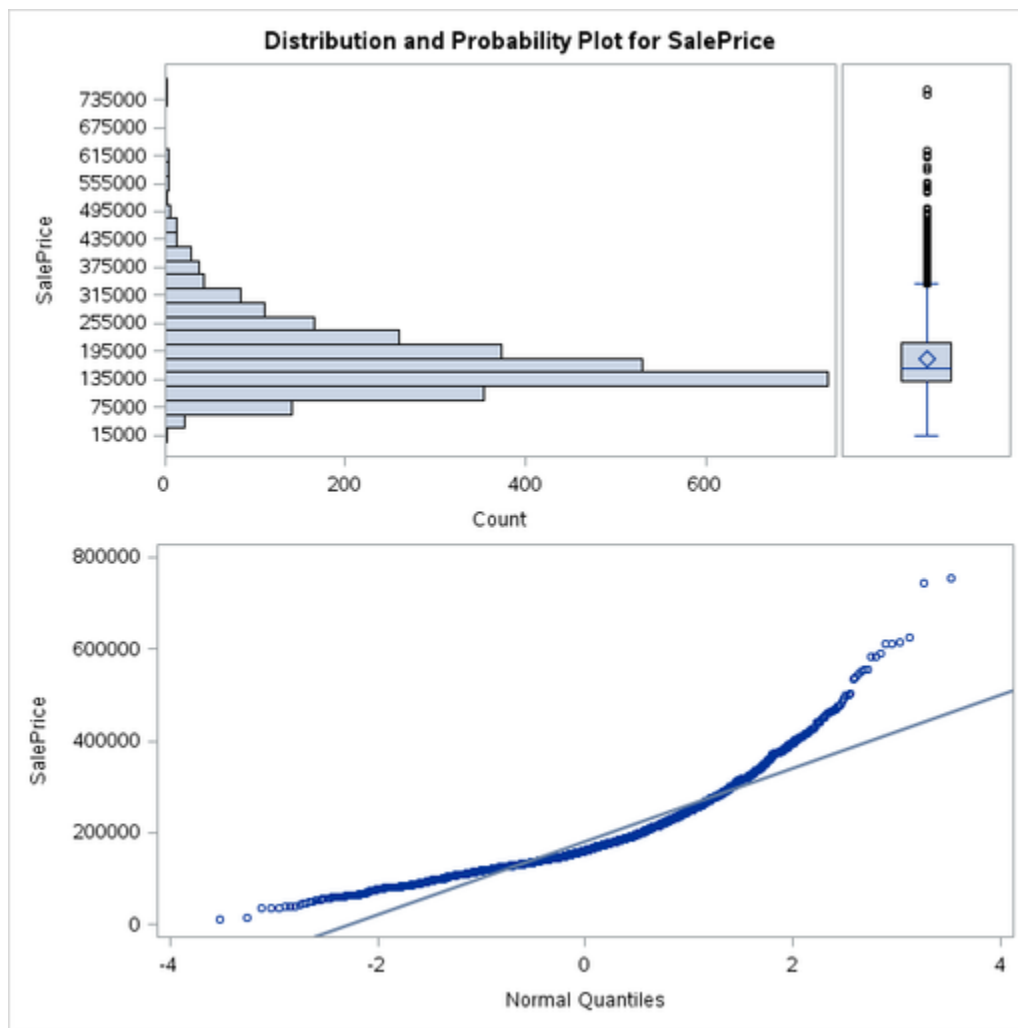


Figure 3: Distribution and Probability Plot for SalePrice

We can see at least a single outlier showing up at the top of the histogram. Looking at the normal probability plot for SalePrice we observe that the line is curved, which suggests a skewed distribution. We also observe some outliers in this depiction of the data.

We'll manually inspect the lowest and highest five observations:

Observation	Value	Type
182	12789	Lowest
1554	13100	Lowest
727	34900	Lowest
2844	35000	Lowest
2881	35311	Lowest
45	611657	Highest
1064	615000	Highest
2446	625000	Highest
1761	745000	Highest
1768	755000	Highest

Table 9: Extreme Observations of SalePrice

We'll also consider the Quantiles:

Level	Quantile
100% Max	755000
99%	457347
95%	335000
90%	281357
75% Q3	213500
50% Median	160000
25% Q1	129500
10%	105250
5%	87500
1%	61500
0% Max	12789

Table 10: Quantiles

If this were a formal analysis we would spend a great deal of time examining a list of the outliers to find out if some of them were clearly *impossible*. In just a quick glance, at the low end, we could see that there was possibility of at least one substantive outlier, and on the top end there was two observations made a few measurements apart that have close to the same value (and both are outliers).

We look at the data (histogram) while laying over a normal distribution. However, we've already stated that it's unlikely that the distribution of observations is normal. In the 'olden' days before we had 256GB of memory (dornick) readily at our fingertips, and decades of statistical software development to tap into, we would consider examining the two standard deviation rule. There are data sets that easily show that this rule isn't applicable to all analyses projects. We'll take a stab at establishing a three bucket outlier processing system using the 99% and 1% quantile values as our high and low band passes, even if it's likely to be a bad mechanism for outlier elimination:

```
* three bucket outlier classification;
data outliers;
  set ames;
  keep SalePrice GrLivArea MasVnrArea BsmtUnfSF;
  if SalePrice <= 61500 then price_outlier = 1;
  else if SalePrice > 61500 & SalePrice < 457347 then price_outlier = 2;
  else if SalePrice >= 457347 then price_outlier = 3;
  keep SalePrice price_outlier GrLivArea MasVnrArea BsmtUnfSF

proc sort data=outliers;
  by price_outlier;

* means of each outlier bucket;
proc means data=outliers;
  by price_outlier;
  var SalePrice;
```

We'll create a new data set by pruning all the outliers at the low and high end of our data set:

```
* prune the outlier data into a new data set;
data pruned;
  set outliers;
  if price_outlier = 1 then delete;
  if price_outlier = 3 then delete;

proc univariate normal plot data=pruned;
  var SalePrice;
  histogram SalePrice / normal (color=red w=5);
```

We will use the univariate procedure again to examine how our data looks now:

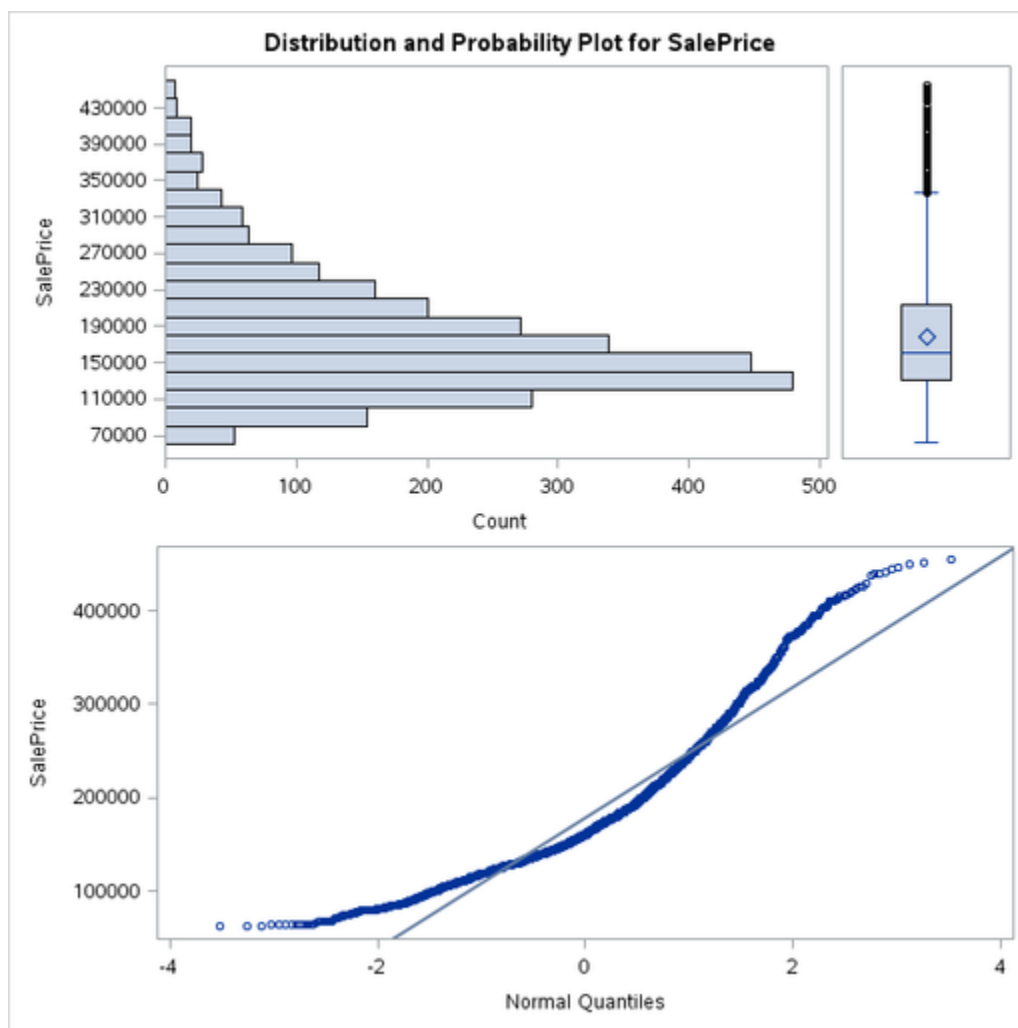


Figure 4: Distribution and Probability Plot for SalePrice (Outlier Pruned)

We now observe that the outliers on both side of this tailed data set have been pruned. The line is still curved, and the histogram is shifted, indicating that we have data that is not necessarily normal in distribution.

From this visual inspection alone we feel little confidence that our outlier identification methodology was sane for this data set. We'll carry forward to examine how pushing the data set farther from normal will impact the model performance.

We'll now manually inspect the lowest and highest five observations again:

Observation	Value	Type
546	62383	Lowest
731	62500	Lowest
744	63000	Lowest
757	63900	Lowest
2602	64000	Lowest

Observation	Value	Type
435	445000	Highest
1027	446261	Highest
1656	450000	Highest
420	451950	Highest
1606	455000	Highest

Table 11: Extreme Observations of SalePrice (Outlier Pruned)

We'll also consider the Quantiles:

Level	Quantile
100% Max	455000
99%	405000
95%	320000
90%	275500
75% Q3	212900
50% Median	160000
25% Q1	130000
10%	107950
5%	91000
1%	75000
0% Max	62383

Table 12: Quantiles (Outlier Pruned)

From this we see that our outlier removal procedure executed correctly, however from comparing the tables and looking at the histograms we don't feel our methodology was sound. It is in our interest to continue with this data set so that we can examine how the performance of linear regression is affected.

Model Comparison on Modified Data Set

We will now compare some simple models with both the original data set and the new data set we've made that removes our classified outliers. The models we're going to compare are:

$$\text{SalePrice} = \beta_0 + \beta_1 \text{GrLivArea} + \epsilon$$

$$\text{SalePrice} = \beta_0 + \beta_1 \text{GrLivArea} + \beta_2 \text{MasVnrArea} + \epsilon$$

$$\text{SalePrice} = \beta_0 + \beta_1 \text{GrLivArea} + \beta_2 \text{MasVnrArea} + \beta_3 \text{BsmtUnfSF} + \epsilon$$

We will use SalePrice as the dependent variable identification for the non-manipulated data set, and SalePrice₀ for the manipulated data set.

After running the regressions, the respective model fits are as follows for the non-manipulated data set:

$$\text{SalePrice} = 13316 + 111.50225 \times \text{GrLivArea}$$

$$\text{SalePrice} = 26596 + 94.57781 \times \text{GrLivArea} + 118.54037 \times \text{MasVnrArea}$$

$$\text{SalePrice} = 25788 + 93.91175 \times \text{GrLivArea} + 118.57390 \times \text{MasVnrArea} + 3.23138 \times \text{BsmtUnfSF}$$

After running the regressions, the respective model fits are as follows for the manipulated data set:

$$\text{SalePrice}_0 = 31604 + 98.20384 \times \text{GrLivArea}$$

$$\text{SalePrice}_0 = 40449 + 85.89009 \times \text{GrLivArea} + 997.44693 \times \text{MasVnrArea}$$

$$\text{SalePrice}_0 = 39321 + 84.92176 \times \text{GrLivArea} + 97.45046 \times \text{MasVnrArea} + 4.61086 \times \text{BsmtUnfSF}$$

The diagnostic model performance information is below:

Model	Adj R-Square	F Value
$\text{SalePrice} = 13316 + 111.50225 \times \text{GrLivArea}$	0.5002	2908.60
$\text{SalePrice} = 26596 + 94.57781 \times \text{GrLivArea} + 118.54037 \times \text{MasVnrArea}$	0.5594	1844.94
$\text{SalePrice} = 25788 + 93.91175 \times \text{GrLivArea} + 118.57390 \times \text{MasVnrArea} + 3.23138 \times \text{BsmtUnfSF}$	0.5595	1231.02
$\text{SalePrice}_0 = 31604 + 98.20384 \times \text{GrLivArea}$	0.4577	2403.02
$\text{SalePrice}_0 = 40449 + 85.89009 \times \text{GrLivArea} + 997.44693 \times \text{MasVnrArea}$	0.5063	1460.07
$\text{SalePrice}_0 = 39321 + 84.92176 \times \text{GrLivArea} + 97.45046 \times \text{MasVnrArea} + 4.61086 \times \text{BsmtUnfSF}$	0.5069	976.10

Table 13: Comparison of model performance with non-manipulated and manipulated data set

We cannot do direct comparison of the Adj R-Square and F Value between the two sets of models, as these calculations are within the context of the model fit to the data set. We can however remark about what the removal of our classified outliers has done to the model performance.

As stated before, we were a bit aggressive with our outlier classification methodology. In this case we can confidently say that the performance of the models was degraded by our outlier removal. If we were to assume that our outlier classification was actually rigorous and based in both domain knowledge, then we may begin to look for the construction of different models.

That is to say, by having two different data sets and comparing regression performance, the practitioner can look at the respective model performance objectively. If you were certain that your ‘cleaned’ data set was more representative of the phenomena you wanted to model, then observing degradation in model performance would tell you that your initial models we’re fitting well to a different phenomena (influenced by the outliers).

Conclusion / Reflection

How does transformation and outlier deletion impact the modeling process and results?

Do these activities benefit or create additional difficulties

Whats the next steps in the modeling process

Straitening Relationships

Power	Name	Comment
2	Square	Try with unimodal distributions that are skewed to the left
1	Raw data	Data with positive and negative values and no bounds are less likely to benefit from re-expression
$\frac{1}{2}$	Square root	Counts often benefit from a square root re-expression
0	Logarithms	Measurements that cannot be negative often benefit from a log re-expression
$-\frac{1}{2}$	Reciprocal square root	An uncommon re-expression, but sometimes useful
-1	Reciprocal	Ratios of two quantities often benefit from a reciprocal

Table 14: The Ladder of Powers

Model Name	x -axis	y -axis	Comment
Exponential	x	$\log(y)$	This model is the “0” power in the ladder approach, useful for values that grow by percentage increases
Logarithmic	$\log(x)$	y	A wide range of x -values, or a scatterplot descending rapidly at the left but leveling off towards the right
Power	$\log(x)$	$\log(y)$	The Goldilocks model: when one of the ladder’s powers is too big and the next is too small, this may be just right

Table 15: Attack of the Logarithms

Procedures

```
title 'Assignment 3';

libname mydata '/scs/crb519/PREDICT_410/SAS_Data/' access=readonly;

* create a temporary variable (data source is read only);
* we will also do some transformations and only keep;;
* the variables that we're interested in for this study.;
data ames;
    set mydata.ames_housing_data;
    log_SalePrice = log(SalePrice);
    log_GrLivArea = log(GrLivArea);
    sqrt_SalePrice = sqrt(SalePrice);
    keep SalePrice log_salePrice sqrt_SalePrice GrLivArea log_GrLivArea MasVnrArea BsmtUnfSF BsmtFinSF1

ods graphics on;

* verify that we did indeed carry over the variables of interest;
proc print data=ames (obs=5);

* lets look at the linearity of the transformations;
proc sgscatter data=ames;
    title 'Normal and Transformed Variables';
    plot (SalePrice log_SalePrice) * (GrLivArea log_GrLivArea);

* comparing models where we do some re-expression;
proc reg data=ames;
    model SalePrice = GrLivArea;
    model SalePrice = log_GrLivArea;
    model log_SalePrice = GrLivArea;
    model log_SalePrice = log_GrLivArea;

* correlation with our re-expressed dependent variable;
proc corr data=ames nosimple rank;
    var GrLivArea MasVnrArea BsmtUnfSF BsmtFinSF1 FirstFlrSF TotalBsmtSF GarageArea;
    with log_saleprice;
run;

proc sgscatter data=ames;
    title 'SalePrice and log_SalePrice vs GrLivArea';
    plot (SalePrice log_SalePrice) * GrLivArea;

proc reg data=ames;
    model SalePrice = GrLivArea;

proc reg data=ames;
    model sqrt_SalePrice = GrLivArea;

* identify outliers in SalePrice;
proc univariate normal plot data=ames;
    var SalePrice;
    histogram SalePrice / normal (color=red w=5);
```

```

* three bucket outlier classification;
data outliers;
    set ames;
    keep SalePrice price_outlier GrLivArea MasVnrArea BsmtUnfSF;
    if SalePrice <= 61500 then price_outlier = 1;
        else if SalePrice > 61500 & SalePrice < 457347 then price_outlier = 2;
        else if SalePrice >= 457347 then price_outlier = 3;

proc sort data=outliers;
    by price_outlier;

* means of each outlier bucket;
proc means data=outliers;
    by price_outlier;
    var SalePrice;

* prune the outlier data into a new data set;
data pruned;
    set outliers;
    if price_outlier = 1 then delete;
        else if price_outlier = 3 then delete;

proc univariate normal plot data=pruned;
    var SalePrice;
    histogram SalePrice / normal (color=red w=5);

* regression model for non-manipulated data set;
proc reg data=ames;
    model SalePrice = GrLivArea;
    model SalePrice = GrLivArea MasVnrArea;
    model SalePrice = GrLivArea MasVnrArea BsmtUnfSF;

* regression model for non-manipulated data set;
proc reg data=pruned;
    model SalePrice = GrLivArea;
    model SalePrice = GrLivArea MasVnrArea;
    model SalePrice = GrLivArea MasVnrArea BsmtUnfSF;

run;

```

References

- [1]W. A. Huber, “In linear regression, when is it appropriate to use the log of an independent variable instead of the actual values?” 2010. [Online]. Available: <http://stats.stackexchange.com/questions/298/in-linear-regression-when-is-it-appropriate-to-use-the-log-of-an-independent-variable>. [Accessed: 17-Apr-2015]
- [2]J. W. Tukey, “Exploratory data analysis,” 1977.
- [3]W. A. Huber, “Logistic regression: Transforming variables,” 2010. [Online]. Available: <http://stats.stackexchange.com/questions/4831/logistic-regression-transforming-variables>. [Accessed: 17-Apr-2015]